Journal of Urban Economics 139 (2024) 103610

Contents lists available at ScienceDirect

Journal of Urban Economics

journal homepage: www.elsevier.com/locate/jue

Volatility in Home Sales and Prices: Supply or Demand?☆

Elliot Anenberg *,1, Daniel Ringo[1]

Board of Governors of the Federal Reserve System, United States of America

ARTICLE INFO

Keywords:
Housing market search
Housing market volatility

ABSTRACT

We use a housing search model and data on individual home listings to decompose short-run fluctuations in home sales and price growth into supply or demand factors, defined as the number of new sellers and buyers entering the housing market, respectively. We find that fluctuations in the number of buyers demanding homes explain much more of the variation in home sales and price growth than do fluctuations in the supply of homes for sale. In our preferred paramaterization, fluctuations in demand explain essentially all of the variation in home sales, and most of the variation in prices. We consider two implications of these results. First, we show that reduction of supply was a minor factor relative to an increased number of buyers in the tightening of housing markets during COVID-19. New for-sale listings would have had to expand 30 percent to keep the rate of price growth at pre-pandemic levels given the pandemic-era surge in demand. Second, we estimate that the number of buyers demanding homes is very sensitive to changes in mortgage rates, even more so than comparable estimates for home sales, suggesting that policies that affect housing demand through mortgage rates can influence housing market dynamics.

1. Introduction

Both home sales and home price growth are volatile and cyclical, generally rising together during booms and falling during busts. These fluctuations have important implications for economic activity, financial stability, and access to homeownership.[2] They are also frequent targets of policy making, as governments have a variety of policy options that can affect the demand for homes (e.g. first-time home buyer tax credits, mortgage subsidies, monetary policy) or the supply of homes available for sale (e.g. zoning reform, tax assessment restrictions, transfer taxes). To predict the outcomes of these policy choices, it is therefore important to understand the determinants of housing market volatility. The main objective of this paper is to estimate the extent to which short-run fluctuations in sales and price growth are driven by the demand for homes, or by the supply of homes for sale.

To decompose these fluctuations into supply or demand factors, we use a simple model in which the stock of active buyers and sellers produce market dynamics through a frictional housing search process. Our housing search model is motivated by Figs. 1 and 2, which show that changes in home sales and prices are accompanied by changes in housing liquidity. Fig. 1 shows that homes take longer to sell (i.e. a low rate of sale hazard) when the stock of active sellers is high, and Fig. 2 shows that the average time on market of for-sale homes is strongly negatively associated with home price growth. These correlations suggest an important role for market tightness, or the ratio of buyers to sellers, in explaining short-run housing dynamics.

We therefore take the number of newly active sellers and buyers entering the market as the fundamental measures of demand and supply we will be investigating. Newly active sellers ("supply") include both builders of new construction and sellers of existing homes, with the latter typically accounting for a large majority of newly active sellers. New buyers ("demand") include both first-time homebuyers and existing homeowners searching for a different home to buy. These inflows of buyers and sellers, balanced with the outflows due to successful matches and discouraged searchers, determine the market tightness in equilibrium.[3]

☆ The analysis and conclusions set forth are those of the authors and do not indicate concurrence by other members of the research staff or the Board of Governors.

* Corresponding author.
E-mail address: elliot.anenberg@frb.gov (E. Anenberg).

[1] Made equal contributIon to the paper.

[2] For home prices and economic activity, see Aladangady (2017), Berger et al. (2017), Guren et al. (2021), Mian and Sufi (2011), Mian et al. (2013). For home sales, see Benmelech et al. (2023), Karahan and Rhee (2019), Ortalo-Magne and Rady (2006).

[3] For ease of exposition, we use this supply and demand terminology throughout the paper. Our notions of supply and demand are not comprehensive measures, however. For example, they do not include shifts in the intensive margin of demand (i.e. conditional on searching for a house, the size or quality of houses buyers are searching for may vary).

https://doi.org/10.1016/j.jue.2023.103610
Received 30 December 2022; Received in revised form 24 October 2023
Available online 2 December 2023
0094-1190/Published by Elsevier Inc.

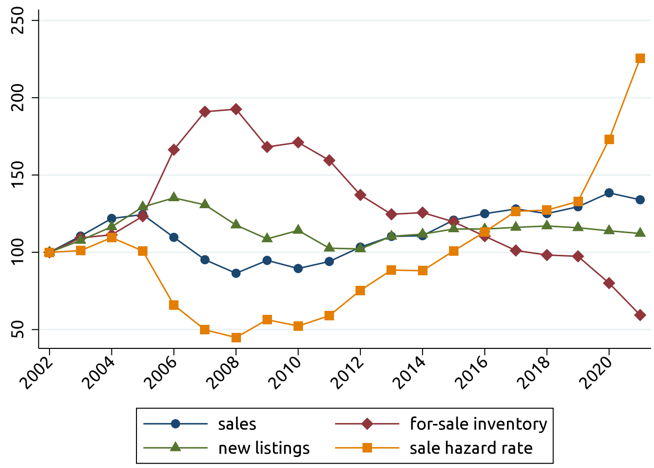

Fig. 1. Annual sales volume, new listings, inventory, and sale hazard rate.
Notes: All series are indexed to 2002 values. The sale hazard rate is calculated as the number of sales contracted each month divided by the number of homes actively listed for sale at some point in the month. The annual sale hazard is the average of the monthly sale hazards, weighted by the number of homes listed for sale each month.

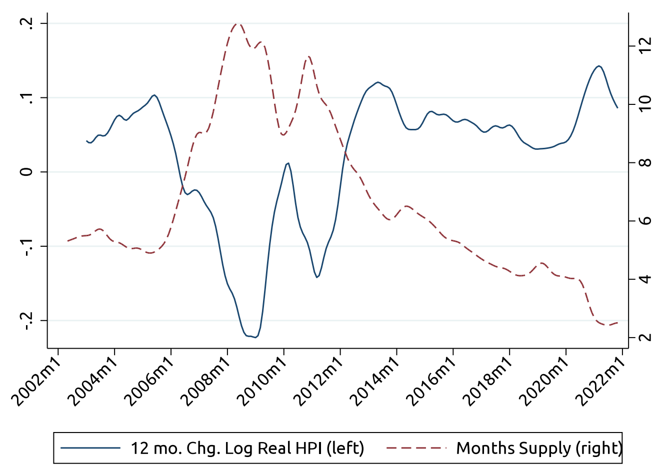

Fig. 2. Months supply and house price growth.
Notes: Months supply is equal to the inverse of the sale hazard rate. House price index shows the estimates of the time dummies from a regression of log house prices on time dummies, home characteristics, and zipcode fixed effects. House price index is adjusted for inflation using the consumer price index excluding shelter.

An important empirical challenge is that we do not have data on the number of potential buyers actively searching for a home. Sellers advertise for-sale homes to buyers through platforms such as the multiple listings service (MLS), but active buyers generally do not record their presence or search activity. Consequently, we obtain data on the inflow of new listings for sale, the stock of active sellers, and the sale hazard rate from the United States MLS, and we use our model structure combined with these MLS data to estimate the number of buyers searching for homes.

Using the model and estimated inflow of buyers, we consider simulations where we hold fixed either the supply or demand sides of the market, and allow the other to vary as in the data. In our preferred parameterization, fluctuations in the number of buyers explain essentially all of the variation in home sales, and most of the variation in prices, between 2002–2021. In other paramaterizations, demand can be forced to play a somewhat less important role, but its strong contribution relative to supply is a robust result. Extending our model to include moving homeowners who enter both sides of the market simultaneously as both buyers and sellers, we show that about half of the variation in sales volume is attributable to the demand of these joint buyer-sellers and half to demand from households that are exclusively buyers. Joint buyer-sellers have very little effect on price growth, however, as their effects on the demand and supply sides cancel out.

We then consider two implications of the results from our model. First, we show that the COVID-19 housing boom in the U.S. was driven by an increase in the number of buyers. Even though the supply of new for-sale listings fell sharply at the beginning of the pandemic, we show that reduction of supply was a minor factor relative to an increased number of buyers in explaining the tightening of housing markets over the first year of the pandemic. A policy concern during the pandemic has been that the sharp rise in house prices has exacerbated affordability pressures and increased financial stability risks. We use our model to estimate how much additional supply would be needed to offset the estimated increase in searching buyers so that house prices continued along their pre-pandemic trend, instead of accelerating. We find that a 30% increase in the monthly number of homes coming on to the market would have been necessary to keep up with the pandemic-era surge in demand. Since new construction typically accounts for about 15% of supply, our estimates imply that new construction would have had to increase by roughly 300% to absorb the pandemic-era surge in demand. This is a very large, unrealistic impulse to housing supply in the short-run, suggesting that policies aimed at reducing bottlenecks to new construction would have done little to cool the housing market during COVID-19.[4]

Second, we show that our estimated number of buyers searching for homes is very mortgage rate elastic. We estimate that a one percentage point increase in the mortgage rate lowers the number of buyers by 10.4 percent. This is a larger demand sensitivity to rates than evidence using purely observable housing market variables suggests. In particular, higher mortgage rates also decrease home sales, but the semi-elasticity is 6, or about one-half the semi-elasticity of the number of buyers. Because search frictions effectively smooth the response of home sales to demand shocks over time, estimates of the short-term elasticity of home sales obscure some of the mortgage rate sensitivity of demand. A high mortgage rate sensitivity of demand combined with our main result showing that short-run housing market fluctuations are largely explained by demand suggest that policies that target mortgage rates are an effective way to influence short-run fluctuations in the housing market.[5]

Our notion of the demand for home buying in this paper is distinct from the aggregate demand for housing services, most of which is fulfilled by households staying out of the market and consuming housing services provided by their current residence. Similarly, the supply of homes for sale is distinct from the aggregate supply of housing services, the quantity of which can be increased by new construction or decreased by demolition or conversion of properties. Both sets of concepts (that of new active market participants, which we use, versus that of aggregate stocks) are important. The literatures on the elasticities of housing demand and supply generally focus on the stock concept.[6] However, housings' contribution to variation in

[4] In the long run, increasing new construction may be a more effective policy response. Longer run fluctuations in the housing market are beyond the scope of this paper.

[5] Our estimates of the rate semi-elasticity of home buying demand is in line with other estimates in the literature using shocks to the cost of credit. Estimated semi-elasticities of demand for home buying range from 20 for lower-income, credit constrained households to 6 for middle- and upper-income households (Bhutta and Ringo, 2021 and Ringo, 2022). Estimates of the elasticity of the *intensive* margin of mortgage demand—that is, the amount borrowed conditional on a loan being taken out at all—are much smaller. For example, DeFusco and Paciorek (2017) estimate the rate semi-elasticity of intensive margin demand to be about 2.

[6] See, for example, Saiz (2010), Gyourko and Molloy (2015), Hanushek and Quigley (1980), Molloy et al. (2022), or Zabel (2004).

national income depends on transactions (as households adjust their housing consumption), rather than passive consumption from the stock, and is therefore tied to the number of active market participants. Furthermore, house prices are largely influenced in the short-run by the balance of active buyers to sellers (i.e. market tightness), as suggested by Fig. 2. While the aggregate supply of housing services cannot adjust quickly, the supply of homes for sale can be quite volatile over the business cycle, as can be seen in Fig. 1. As we discuss further in Section 6, volatility in supply can be accompanied by volatility in market tightness because many sellers are not also buyers.

Our paper is related to a large literature that takes a search-theoretic approach to modeling the dynamics of the housing market. Han and Strange (2015) provide a summary of this literature. Our model of random housing search is similar to a number of models in the literature – see, for example, Diaz and Jerez (2013), Guren and McQuade (2020), Krainer (2001), Novy-Marx (2009), Piazzesi and Schneider (2009), and Gabrovski and Ortego-Marti (2019). Our contribution, building on Anenberg and Ringo (2021), is to show that this simple model of housing search combined with time-series data on new listings, for-sale inventory, and withdrawals can be used to estimate the demand for homes.[7] We use this estimate of buyer demand along with our data and model to provide new results on the contribution of supply and demand to volatility in the housing market.

Our paper is closely related to the empirical work in Ngai and Sheedy (2020) (henceforth NS). Using data from the U.S., NS find that variation in the supply of homes for sale explains essentially all of the volatility in sales volume, whereas we find that supply has very little explanatory power for volatility in sales volume. Motivated by their empirical finding that sales volume dynamics are explained by supply, NS also build a search-and-matching model with endogenous house prices and moving decisions.

In Section 7, we show that our results differ from the empirical results in NS for two reasons. First, we make different assumptions about how supply affects the sale hazard rate. In NS, the stock of active listings for sale does not affect the sale hazard rate. In our model the sale hazard depends on the market tightness, meaning it varies endogenously with the supply of homes for sale. As supply increases (all else equal), the sale hazard goes down and offsets much of the effect of increased supply on sales volume. Our results highlight the importance of modeling market tightness, and its implications for the matching process, when evaluating the relative roles of demand and supply. Second, we use micro data on individual listings that allow us to directly measure supply – i.e. the inflow of new listings for sale. NS use aggregate data and as a result of data limitations, their measure of supply is actually new listings *net of withdrawals*. Because withdrawals are negatively correlated with demand, the NS measure of supply is influenced by demand factors, causing their estimates to overstate the contribution of supply to volatility in sales volumes.[8]

Our finding that demand is important for explaining variation in sales volume is consistent with and related to a number of recent papers providing quasi-experimental evidence that sales volumes are sensitive to demand stimulus. Bhutta and Ringo (2021) and Anenberg and Ringo (2022) find that changes in mortgage rates have important effects on home sales. Berger et al. (2020) find that a national first-time homebuyer tax credit, a similarly demand-side policy, had a meaningful stimulative impact on home sales. Best and Kleven (2017) find that sales volumes in the U.K. are sensitive to transaction taxes. However, while the statutory incidence of this tax falls on the buyer, their paper does not attempt to determine whether the change in volume occurred through a demand or supply response. We are not aware of any quasi-experimental studies of the effect of the stock of for-sale listings on sales volumes.

2. Data and motivating empirical patterns

Our data are MLS records provided by CoreLogic. The data come directly from regional boards of realtors, and cover over 50 percent of property listings in the U.S. Information on homes listed for sale includes the initial listing date, the withdrawal date if the home is removed from the MLS without a sale and is not subsequently relisted a short time later, the contract date if the home is sold to a buyer, the asking price, and many home characteristics, including the address.[9] The MLS data have some advantages for our purposes over aggregated listings data, such as those published by the National Association of Realtors (NAR). The data on individual listings allow us to observe the actual inflow of new listings as opposed to inferring it from net changes in total for-sale listings and home sales, a procedure that can lead to mismeasurement caused by withdrawals and homes that list and sell within the same month. Furthermore, the listing-specific data allows us to control for characteristics of the house or listing that could affect the sale hazard (e.g. compositional effects or whether the seller has set an asking price well above or below prevailing prices).

The data run from 2002–2021. From the full sample, we select a subset of 263 counties due to data limitations for some counties. We describe these limitations as well as our procedure for selecting counties in the Appendix.

Fig. 1 shows trends in sales volume, new listings, for-sale inventory, and the sale hazard rate over our sample period. The sale hazard rate is calculated as the number of sales contracted each month divided by the number of homes actively listed at some point during the month. The annual sale hazard is the average of the monthly sale hazards, weighted by the number of homes listed for sale each month. Sales volume rises during the early 2000s, and then falls sharply during the Great Recession. Sales volume slowly recovers from its fall and only in recent years has the level of sales volume returned to early 2000s levels.

One might expect sales volume to be closely related to new listings, as homes can only transact if they are put on the market for sale. Remarkably, however, Fig. 1 shows that new listings are only weakly correlated with sales volume over our sample period. New listings and home sales both rise during the early 2000s, but then diverge as sales volume declines during the Great Recession and new listings remain elevated. New listings have remained fairly flat over the last decade even as sales volume has been on a strong upward trend. These trends suggest that understanding the behavior of new listings alone is not sufficient for understanding cyclicality in sales volume.

In contrast, the figure shows that sales volume and the sale hazard rate have a very high correlation. In addition, Fig. 2 shows that the sale hazard is also strongly associated with house price growth. The figure shows the 12-month change in a real house price index and the “months’ supply”.[10] Months’ supply is the ratio between the number of homes for sale and the number of sales, or the inverse of the monthly sale hazard. Months’ supply alone can explain about 80 percent of the variation in house price growth over our sample period. The strong associations of the sale hazard rate with sales volume and house price growth suggest that understanding the drivers of the sale hazard rate

[7] Concurrent work by Gabrovski and Ortego-Marti (2021) uses a similar model to estimate the pool of active buyers in order to estimate the slope of the Beveridge curve in the housing market. To estimate the pool of active buyers, Gabrovski and Ortego-Marti (2021) use Census data on vacancies and time-to-sell for new home sales.

[8] Withdrawals are negatively correlated with demand because sellers often become discouraged after failing to find a buyer for an extended period of time. For additional evidence on the counter-cyclicality of withdrawals, see Carrillo and Williams (2019).

[9] CoreLogic treats properties that are withdrawn from the market and then relisted a within two months as a single, continuous listing.

[10] We compute the quality-adjusted house price index from our micro data using a standard hedonic price regression, as described in the Appendix.

Table 1
County-level Growth in sales, new listings and sale hazard.

	Sales growth	Sales growth	Sales growth	Sales growth	Sales growth	Sales growth
New Listings Growth	0.346***		0.330***		0.244***	
	(0.0417)		(0.0443)		(0.0346)	
Sale Hazard Growth		0.732***		0.731***		0.835***
		(0.0293)		(0.0290)		(0.0188)
County FE			X	X	X	X
Month-Year FE					X	X
R-squared	0.0827	0.695	0.0955	0.704	0.192	0.788
N	54 026	54 026	54 026	54 026	54 026	54026

Standard errors in parentheses
* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.
Note: OLS regression results of the 12-month growth in sales volume on the 12-month growth in new listings and the sale hazard rate. Each variable is measured at the county-month level and each regression pools observations across counties. Standard errors two-way clustered by county and month.

is key for understanding the cyclicality of the housing market. The importance of the sale hazard rate motivates our model of housing search in Section 3. Our model focuses on this sale hazard rate and allows us to predict how the sale hazard rate changes with supply and demand.

2.1. County-level evidence

County-level statistics provide additional motivating evidence for the limited role new listings play in explaining short-run variation in sales volume. Table 1 reports regression results of the 12-month growth in sales volume on the 12-month growth in new listings and the sale hazard rate. Each variable is measured at the county level and each regression pools observations across counties. The first column shows that a one percentage point increase in new listings growth is associated with a 0.35 percentage point increase in sales volume growth, but the R^2 is only 0.08. New listings growth alone explains a tiny fraction of the variation in sales volume growth. The second column shows that a one percentage point increase in sale hazard growth is associated with a 0.73 percentage point increase in sales volume growth, and the R^2 from this regression is 0.70. Growth in the sale hazard explains almost ten times more of the variation in sales volume growth than does new listings growth. The remaining columns show that the results are similar when fixed effects are added for county, or for county and year-month.

3. Model of housing search

To facilitate our decomposition of housing market cyclicality into demand and supply factors, we use a simple model of housing search. We define supply as the *flow* of homes coming onto the for-sale market each period. Similarly, we define demand as the *flow* of prospective buyers that enter the market to find a home to purchase. Our measures of supply and demand are thus based on the decision to participate in the housing search market. Our model abstracts from heterogeneity in the housing stock and decisions of buyers to search for smaller or larger homes. In our model, supply and demand affect the housing market equilibrium through their effect on market tightness, θ, which in turn affects the rate at which homes are sold. Market tightness is the ratio of the *stock* of prospective buyers, b, relative to the stock of sellers, s: i.e. $\theta = \frac{b}{s}$.

The stock of sellers (i.e. the inventory of homes for sale), s, in each month t evolves as

$$s_{t+1} = s_t - s_t q_t^s(\theta_t) - s_t(1 - q_t^s(\theta_t))w^s + n_{t+1}^s \tag{1}$$

where q^s is the rate at which homes are sold, w^s is the rate at which unsold homes are withdrawn from the market, and n^s is the inflow of new sellers (i.e. our fundamental measure of supply). Eq. (1) expresses the stock next period as the stock this period (first term) minus the outflow arising from sales and withdrawals (middle terms) plus the inflow (final term).

Similarly to the supply of homes for sale, there is a stock of currently-searching buyers, b, that is replenished by an inflow of new buyers, and depleted as buyers purchase a home and exit the market, or drop out without purchasing. The stock of buyers evolves as

$$b_{t+1} = b_t - b_t q_t^b(\theta_t) - b_t(1 - q_t^b(\theta_t))w^b + n_{t+1}^b \tag{2}$$

where q^b is the rate at which a buyer finds a home to buy, w^b is the rate at which buyers leave the market, and n^b is the inflow of new buyers.

Buyers and sellers interact via the search and matching process, which we model as Cobb–Douglas with constant returns to scale. Under Cobb–Douglas, the probabilities of buying and selling are:

$$q_t^s(\theta_t) = \theta_t q_t^b(\theta_t) = A_t \theta_t{}^\eta \tag{3}$$

where $0 < \eta < 1$ is the elasticity of the probability of sale with respect to market tightness and A is a parameter that determines the efficiency of the matching function. We allow A_t to vary over time based on factors exogenous to our model, discussed further in Section 4.1. We use the Cobb–Douglas function because the calibration of η makes explicit our assumption about the marginal effects of s and b. As described in Section 4.2 we can calibrate this parameter to quasi-experimental estimates in the literature, and in Section 5.1.1 test the extent which our results depend on a particular parameterization. We further discuss this choice of functional form and show robustness to alternative specifications of the search-and-matching function in the Appendix.

Because q^s is an increasing function of market tightness, the more sellers there are in the market, the slower a given house is likely to sell (all else equal). Sellers crowd each other out and create congestion by competing for the stock of available buyers — the more houses there are for sale, the less likely any particular house is to receive an offer. This prediction is consistent with the very strong negative correlation in the data between the number of sellers on the market, s, and the sale hazard rate, q^s shown in Fig. 1.

An alternative to the random search model described above would be to model housing search via stock-flow matching. For example, Smith (2020) uses such a stock-flow model with endogenous seller entry to explain hot and cold housing markets. In the Appendix, we show robustness of our main results in a stock-flow model where homes that have just come to market are more efficient searchers than homes that have been on the for-sale market for some time. Another possible modeling choice would be directed search. Albrecht et al. (2016) develop a directed search model where motivated sellers choose low list prices and relaxed sellers choose high list prices, leading to shorter and longer time-to-sell, respectively.[11] Our model abstracts from the list

[11] List prices play a similar role in the quantitative directed search model in Hedlund (2016).

price decision, but, as we describe below in Section 4.1, we incorporate list prices by allowing them to affect matching efficiency. As a result, a change in the composition of sellers–for example, from relaxed, high list-price sellers to motivated, low list-price sellers–can affect the sale hazard rate without affecting our estimate of the demand for homes.

4. Estimation and calibration

To impute the relative influence of supply and demand on housing market fluctuations, we need estimates of n_t^b and n_t^s. As discussed above, n_t^s is directly observed from our listings data. We seasonally adjust it using the Census Bureau's X-12-ARIMA seasonal-adjustment program. We estimate n^b using our model structure and the listings data. Our approach to estimating n^b requires estimates of the sale hazard, q^s; estimates of the matching efficiency, A; and calibration of several parameters. We discuss each in turn.

4.1. Estimating sale hazard, q^s, and matching efficiency, A

We estimate sale hazards and matching efficiency using our panel of active listings at a monthly frequency. Houses enter the panel either in the month they are listed for sale, or in January 2002 if the listing was already active at that point. They exit when the house is delisted, and the panel as a whole ends in November 2021. Some homes are delisted because a sale has occurred, others are delisted because the seller has decided to no longer market the home for sale. Homes that are delisted from the market without a sale are treated as censored observations.

Using this sample, we estimate a time period specific sale hazard for each month of the panel. This sale hazard is intended to represent that of a generic listing, affected only by the number of active buyers and sellers, so we need to control for variation in the composition of listings that could affect sale hazards. To accomplish this, we estimate an accelerated failure time model where the hazard rate of sale for house i at time t is parameterized as

$$h_{it} = exp(\delta_t + \beta^A X_{it}^A) \tag{4}$$

where δ_t denotes a set of month-year fixed effects and X^A is a vector of observables that affect matching efficiency. In X^A, we include characteristics of the home, such as its age and number of bathrooms, as well as the home's list price relative to an expected market sales price.[12] These listing-specific characteristics could affect sale hazards for reasons external to the count-based notions of supply and demand we are concerned with. For example, very old homes may not be suitable matches for many buyers. If the pool of homes for sale happen to be older than is typical, matching efficiency for that time period could be low. A high list price relative to market sales price could proxy for a low seller search intensity or unrealistic expectations, also lowering matching efficiency. A mismatch between seller expectations and buyer willingness to pay was especially relevant during the years of the financial crisis, when homeowners were slow to accept how much the price of their homes had fallen.

The effect of X^A on probability of sale is identified using cross-sectional variation, and the month-year fixed effects (δ_t) capture residual variation in average sale hazard over time that is not related to X^A. We interpret variation in δ_t as variation in the sale hazard rate that is related to variation in market tightness. We use an exponential hazard function because Eq. (3) implies that A_t has a proportional effect on the sale hazard given the market tightness. In Eq. (4), the log of the sale hazard is additively separable in the logs of matching efficiency and market tightness, consistent with Eq. (3).

[12] The expected market sales price is the predicted value from a auxiliary regression of log sales prices on home characteristics and month-year dummy variables.

Our estimate of the sale hazard rate is just the average predicted value from estimating Eq. (4):

$$\hat{q}_t^s = \frac{1}{N_t} \sum_{i=1}^{N_t} exp(\hat{\delta}_t + \hat{\beta}^A X_{it}^A) \tag{5}$$

where N_t denotes the number of homes listed on the market in period t.

We estimate A_t as

$$\hat{A}_t = \bar{A} \frac{1}{N_t} \sum_{i=1}^{N_t} exp(\hat{\beta}^A X_{it}^A) \tag{6}$$

which is the estimated sale hazard rate net of the estimated contribution of the month-year fixed effects. Because the month-year fixed effects can only be estimated relative to a baseline period, A_t is identified up to a scale parameter, $\bar{A}$.

We also estimated a much simpler alternative specification which attributes all variation over time in the sale hazard to market tightness (i.e. do not allow the sale hazard to vary with X^A). This simpler specification where A_t is fixed in every period yields qualitatively similar results to our main specification for A_t described in this subsection.

4.2. Calibration of parameters

As is common in the housing literature with random search, we calibrate the elasticity of the matching function, η, to 0.84 to match the estimate from Genesove and Han (2012).[13] Genesove and Han (2012) estimate η using cross-market regressions and survey data on buyer time-on-market, seller time-on market, and number of homes visited by buyers. Subsequent work has validated this estimate using a variety of different identification strategies. Head et al. (2014) calibrate η in their housing search model to target the relative volatility of sales growth to income growth, and arrive at $\eta = 0.86$. Recent work by Grindaker et al. (2021) also arrives at almost an identical estimate to Genesove and Han (2012) using a shift-share shock to market tightness.

Under our matching function, the addition of an extra buyer or seller to the market increases sales volume, but not one-for-one, as the addition of an extra buyer (seller) creates competition or crowd out for other buyers (sellers), lowering the probability of a match. The calibration of $\eta = 0.84$ implies that the addition of an extra buyer to the market has a relatively low crowd out effect on the probability that other buyers in the market match with a for-sale home. Adding an extra seller to the market, however, has a comparatively larger (negative) effect on the probability that other sellers in the market match with a buyer.[14] Genesove and Han (2012) discuss how these crowd-out results could be generated by the MLS. The MLS allows buyers to observe all for-sale listings, but sellers cannot typically observe anything about the pool of potential buyers or take active steps to match with a particular buyer. As a result, buyers can more easily and quickly substitute to other listings when multiple buyers are interested in the same house (i.e. if it sells just before they tour it). Sellers, in contrast, must passively wait for interested buyers to arrive.

We calibrate $\bar{A} = 1.4$ using survey data from the NAR on average search time for buyers in 2019.[15] Buyers reported searching for 10 weeks on average, and we calibrate $\bar{A}$ so that the median buyer simulated in 2019 matches in this time frame.

[13] For example, Anenberg and Bayer (2020), Guren and McQuade (2020), Guren (2018) also calibrate to Genesove and Han (2012).

[14] To see this, note that $\partial M/\partial b = A\eta\theta^{\eta-1}$ and $\partial M/\partial s = A(1-\eta)\theta^{\eta}$ where $M = sq^s$ denotes the number of matchings or sales. The relative crowd out effects depend on θ, but except in very tight markets (i.e. those with very large values of θ), $\partial M/\partial b > \partial M/\partial s$ for $\eta = 0.84$. When market tightness is high, the addition of an extra seller does relatively more to stimulate sales than when tightness is low.

[15] Source: National Assocation of Realtors (2019)

For our counterfactual simulations, described below in Section 5, we calibrate $w^s = 0.061$ to match the average monthly withdrawal hazard in our MLS data. In those simulations, we hold the withdrawal hazard fixed to ensure we are isolating the variation in housing market outcomes due exclusively to supply or demand, respectively. As described in Section 7, the withdrawal hazard of for-sale listings is a function of demand conditions, and failing to distinguish changes in the inflow of new listings from changes in the outflow of withdrawals can lead to spurious inference. We do not have any data or external estimates to inform w^b, and the dynamics of the model depend on the *net* inflow of potential buyers less withdrawals, rather than the gross inflows and withdrawal outflows individually. Therefore, for simplicity we normalize $w^b = w^s$. Our estimates of the inflow of new buyers, n^b, as described below can therefore be thought of as the fluctuations over time in the net inflow of new buyers less withdrawals, up to a constant scalar determined by our normalization of average w^b.

4.3. Inferring demand, n^b

In this section, we present the equation that expresses n^b in terms of variables and parameters that we can observe or estimate using our data. First, note that by inverting Eq. (3), we can express the number of buyers in any period, which is unobserved in our data, as

$$b_t = s_t\left(\frac{q^s_t}{A_t}\right)^{\frac{1}{\eta}} \tag{7}$$

This equation provides an estimate of b_t because s_t is observed in our data, η is a parameter that we calibrate, and Eqs. (5) and (6) give estimates of q^s and A. The estimate of b_t is essentially a residual and so we would attribute any unobserved shocks that affect matching efficiency to variation in b_t. As described in Section 4.1, however, observable characteristics of listings that we can control for in X^A do very little to explain the time-series variation in sale hazards. Notably, even deviations between initial list price and expected sale price (which can be considered a measure of how motivated sellers are) have little explanatory power. We therefore believe that variation in sale hazards, conditional on the number of listings and their observed characteristics, is largely due to variation in demand and that our estimates of b_t are a good proxy for the number of actively searching buyers.[16] Fig. 3 plots our estimates of the number of buyers (b_t) along with the observed number of sellers (s_t).

Second, by plugging Eq. (3) into Eq. (2) and rearranging, we can express n^b as

$$n^b_t = b_t - b_{t-1} + s_{t-1}q^s_{t-1} + b_{t-1}\left(1 - \frac{s_{t-1}q^s_{t-1}}{b_{t-1}}\right)w^b \tag{8}$$

Given Eq. (7), the right-hand side of Eq. (8) depends only on variables that can be observed or estimated, allowing us to estimate n^b_t. Prior to estimating n^b_t, we seasonally adjust s_t, q_t, and A_t using the Census Bureau's X-12-ARIMA seasonal-adjustment program.

5. Partialling out the effects of demand and supply

Our next step is to infer how much of the variation in housing market activity is driven by n^s_t and $\hat{n}^b_t$, observed supply and estimated demand, respectively.

[16] A potential confounder we cannot easily control for is changing technology. If matching has grown more efficient over time, due e.g. to more widespread adoption of online listings platforms, that trend would introduce a spurious upward time trend in our estimates of b_t. However, Fig. 1 does not indicate a general rising trend in sale hazards over the past 20 years. Instead, we see variation in sale hazards is dominated by the business cycle (the boom-bust-recovery surrounding the global financial crisis followed by the extraordinary circumstances of the COVID-19 pandemic).

Fig. 3. Number of active buyers and sellers.
Notes: Number of buyers is estimated used Eq. (7). Sellers is the number of homes actively listed for-sale in the CoreLogic MLS data.

Our first counterfactual scenario isolates the effect of supply by holding demand constant and letting supply follow its observed course. That is, we set the simulated inflow of potential buyers $\widetilde{n^t_b}$ equal to the estimated sample mean, $\bar{n}_b$, in every period. The simulated inflow of new listings, $\widetilde{n^t_s}$, is set equal to the observed n^t_s. We initialize the market using the observed values of the state variables in January 2002, and simulate forward following Eqs. (1), (2), and (3) of our model. Our second scenario isolates the effect of demand by allowing demand to follow its estimated course, while holding supply constant. We set the simulated inflow of buyers equal to its estimated monthly values, $\widetilde{n^t_b} = \hat{n}^t_b$. The simulated inflow of new listings is instead set to equal its sample mean ($\widetilde{n^t_s} = \bar{n}_s$). Again, the market is initialized at January 2002 levels and simulated forward. In both scenarios, A_t, which captures changes in observable characteristics of listings that affect the sale hazard, is held fixed at its sample mean. Withdrawal hazards for both buyers and sellers, w^b and w^s, are likewise fixed at a constant in both simulations to isolate the effect of new entrants.

We evaluate the fit of each simulation to the data using the formula $R^2 = 1 - \frac{sd(\hat{x}_t - x_t)}{sd(x_t)}$ where x is the true data, $\hat{x}$ is the predicted data under a particular simulation, and sd denotes the standard deviation function.

Two remarks on our simulations are in order. First, in reality supply and demand can move independently as in our simulations due to the large presence of buyer-only or seller-only households such as first-time buyers or households leaving the owner-occupied housing market. However, approximately half of the market consists of households attempting to move from one owner-occupied home to another, entering the market as both buyers and sellers.[17] For this reason, fixing the entirety of either the supply or demand side constant while letting the other vary fully may not be a realistic economic counterfactual for the housing market. Still, these joint buyer-sellers make choices about their dual-search problem, including whether to enter the market first as a buyer, first as a seller, or as a buyer and seller simultaneously. Our simulations help us understand which side of the market their potential entry in matters more. In Section 6 we discuss these joint buyer-seller households, estimate their presence in the housing market over time, and augment our analysis to explicitly account for them.

Second, our simulations treat households' entry into the market as exogenous to prevailing market conditions. In reality, many households' entry decisions will be endogenous to market tightness and house prices. Consequently, our counterfactual simulations should not

[17] See DeFusco et al. (2022), Anenberg and Ringo (2022), and the calculations in the Appendix of this paper.

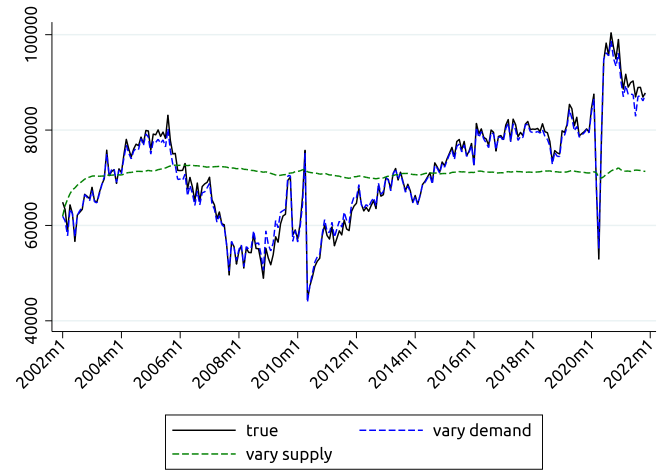

Fig. 4. Sales volumes, observed and counterfactual.
Notes: "True" is the actual sales volume in the data. "Vary supply" is the counterfactual sales volume according to our model when demand is held fixed at its sample mean, but supply varies as in the data. "Vary demand" is the counterfactual sales volume according to our model when supply is held fixed at its sample mean, but demand varies as in our estimates.

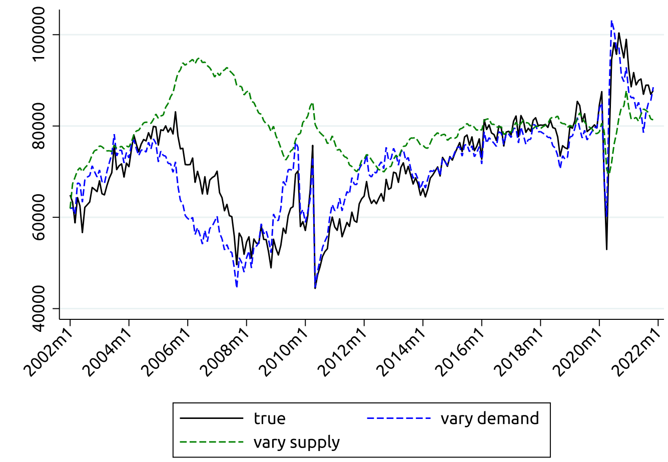

Fig. 5. Sales volumes, alternative crowd-Out.
Notes: Shows simulated sales volume for an alternative calibration of the elasticity of the matching function: $\eta = 0.16$. "True" is the actual sales volume in the data. "Vary supply" is the counterfactual sales volume according to this model when demand is held fixed at its sample mean, but supply varies as in the data. "Vary demand" is the counterfactual sales volume according to this model when supply is held fixed at its sample mean, but demand varies as in our estimates.

be interpreted as predictions of the likely outcome if a central planner were able to eliminate all variation from one side of the market or other. Rather, they inform us about the relative importance of typical fluctuations in the entry of marginal buyers or sellers. This information is necessary to understand the effects of policies that affect the flows of market participants. Furthermore, the simulations help us understand what forces lie behind given market conditions: when housing markets boom (or bust) where should we look for explanations?

5.1. Sales volume

Fig. 4 displays the monthly volume of home sales recorded in our MLS data, alongside counterfactual values from the constant-demand and constant-supply simulations described above. The simulated volumes with constant demand and time-varying supply are nearly flat and only weakly correlated with observed sales. The implication is that variation in supply has essentially no explanatory power over sales volumes.

In contrast, the simulated volumes with constant supply but time-varying demand match the realized sales data very well. There is some small deviation between the series during the Great Recession period, when sellers were particularly likely to set asking prices well above what buyers were willing to pay, but variation in demand explains the overwhelming majority of variation in sales.

5.1.1. Model choice and crowd-out

An important parameter in our model is the elasticity of the matching function, calibrated to $\eta = 0.84$. The closer to 1 is η, the more crowd-out each listing creates. That is, a marginal new listing is more likely to poach a buyer from a different listing, rather than create a new sale. Marginal potential buyers, on the other hand, create little crowd out and so overall sales volume should be quite responsive to variation in the number of buyers. Variation in supply could thus be more important for smaller values of η.

As discussed above, the empirical literature without exception finds evidence for a high value of η. To gauge the sensitivity of our results to lower values of η, however, we show results for $\eta = 0.16$. This alternative calibration is symmetric to the baseline calibration, except buyers, and not sellers, create more crowd out. Results from these alternative simulations are shown in Fig. 5. The reduced supply-side crowd out is clearly visible in the greater variation in sales generated by the constant-demand, varying-supply simulations. However, this simulation still makes a poor fit to the actual data, with a negative R^2. The varying-demand, constant-supply simulations, on the other hand, still do much better in matching the data. While the fit is not as good as in our preferred specification, these simulated sales fit the actual sales volume data with an R^2 of 0.53.

The impotence of supply (and hence dominance of demand) follows from the data and is essentially a requirement of the observation that total for-sale listings and sales volumes are not well correlated. As can be seen in Fig. 1, over the time period we study the correlation is actually substantially negative. If the supply of homes for sale were the major determinant of sales volumes, total listings and sales should generally move together. Under a calibration with more buyer than seller crowd out, the logic of the search model implies an even greater pro-cyclicality of demand to match the observed negative correlation between listings and sales than it does in our preferred calibration. External evidence that crowd-out is actually greater on the supply than on the demand side (motivating a high value of η) only reinforces the relative importance of demand over supply.

5.2. Months' supply and house prices

With the various counterfactual series of b and s simulated as described above, we can also simulate counterfactual paths of the months' supply of homes for sale, shown in Fig. 6. Compared to sales volumes, months' supply is more responsive to variation in new listings. While variation in new listings has little effect on the denominator of the months' supply ratio (sales volume), it is an important determinant of the numerator (total active listings). The figure shows that the vary-supply counterfactual rises and falls between 2002 and 2012, which is consistent with the general pattern in the data over this time period. However, the counterfactual simulation shows an increase in months' supply from 2012 onward whereas in the data, months' supply declined over this time period. Overall, the vary-supply simulated months' supply fits the true months' supply data with an R^2 of 0.13.

As with the sales volume counterfactuals, we find that the vary-demand simulation explains a larger share of the variation. The vary-demand simulation follows the true path of months' supply closely, though it cannot account for the full rise in months' supply during the Great Recession. The elevated level of new listings over this time period contributed significantly to the rise in months' supply. The R^2

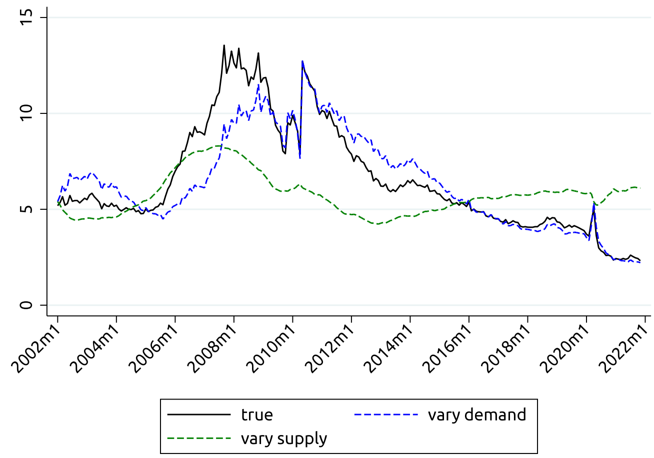

Fig. 6. Months' supply, observed and counterfactual.
Notes: Months' supply is equal to the inverse of the monthly sale hazard rate. "True" is the actual months supply in the data. "Vary supply" is the counterfactual months supply according to our model when demand is held fixed at its sample mean, but supply varies as in the data. "Vary demand" is the counterfactual months supply according to our model when supply is held fixed at its sample mean, but demand varies as in our estimates.

for the vary-demand simulation is 0.55, substantially higher than the vary-supply R^2.

Months' supply of homes for sale is a figure of particular interest because of its close connection with house prices. As shown in Fig. 2, there is a very tight negative correlation between months' supply and house price growth, a relationship sometimes referred to as the housing Phillips curve. The R^2 of a regression of house price growth on months' supply is about 0.80. Caplin and Leahy (2011) and Guren (2018) discuss how a negative relationship between the *level* of months' supply and *changes* in house prices is difficult to explain in a model with full information and rationality. While our housing search model does not take a stand on the specific house price formation process, one way to motivate this tight connection is with a model in which sellers have only limited information about the demand for their homes, but can observe the recent experience of other for-sale listings. A tight market, as evidenced by the rapid sale of recent listings, informs sellers that they can raise prices. Slower sales, indicating that the market is not so tight, would inform them that they may need to lower prices to make a timely sale. Price adjustments to rebalance the number of prospective buyers and sellers willing to transact at current price levels (i.e. market tightness) cause the negative correlation between price growth and months' supply. We describe such a model in the Appendix. Similar intuition for the housing Phillips curve can be found in the models of Carrillo et al. (2015) and Guren (2018).[18]

Taking the reduced-form relationship in Fig. 2 as given and feeding counterfactual months' supply from Fig. 6 into that relationship, we find that price fluctuations in the short term mostly depend on the number of buyers. Variation in supply also has a meaningful influence on months' supply and hence price growth, though its influence is much smaller.[19]

This inference about the determinants of price growth assumes that the observed relationship between market tightness and prices is causal, or mechanically linked as in the model described in the Appendix. However, it is possible that the effect of a shock to market tightness on price growth would not be as strong as the tight correlation apparent in Fig. 2 might suggest. Prices and months' supply could both be influenced independently by some third factor. Lacking a clean source of quasi-experimental variation in demand and supply, we cannot be certain that the true effect is as strong as the observed correlation. Nonetheless, given the clear theoretical connection and that tight correlation, it seems very likely that market tightness (and, consequently, demand) is the primary driver of short-run house price growth.

5.3. Explaining cross-sectional variation in housing markets

The analysis shown thus far has been exploiting and explaining time-series variation in the U.S. housing market in the aggregate. Housing market dynamics differ widely by locality within the country, however. In the Appendix we show that, across counties, variation in sales growth can also be well explained by variation in demand but not at all by variation in supply. Similarly to the results from the aggregate time series, this inference does not rely on a calibration in which there is more crowd out on the supply side than on the demand side (although it is reinforced by such a calibration). As in the aggregate results, supply plays a larger role in explaining cross-sectional variation in months' supply than in explaining sales volumes.

6. The role of joint buyer-sellers

In Section 5, we distinguished the relative importance of supply and demand by running counterfactual simulations in which one series varies and the other is held constant. Economic agents, as described in the model in Section 3, enter exclusively as either buyers or sellers and the inflow of both types can be manipulated independently. As mentioned previously, this dichotimization well describes about half of the housing market. On the demand side, first time-home buyers, vacation home buyers, and investors can all enter the market to buy without putting a home up for sale. On the supply side, other investors or households exiting owner-occupancy (into rentership or group quarters, or by combining households, emigrating or dying) can enter the market as sellers without searching for another home to buy.

The other half of home sales, however, are conducted by joint buyer-sellers. These are owner-occupant households who intend to move, buying their new home and selling their old home nearly simultaneously. Adding such a household to the pool of buyers would necessitate adding them to the pool of sellers as well (and vice versa), meaning their supply and demand cannot be manipulated independently as in our simulations above. A consideration of the distinct contribution of joint buyer-sellers is important because they comprise a large share of the market and the time series of their entry may be different from the entry of buyer-only or seller-only households. In addition, joint buyer-sellers may respond to policy differently than buyer-only or seller-only households, so it is important to understand whether joint buyer-sellers are an important driver of housing market volatility.[20]

[18] In Carrillo et al. (2015), sellers are slow to realize when there has been a shock to the number of buyers searching on the market, and that their bargaining power has consequently changed. Market tightness improves the seller's relative bargaining position and hence increases sales prices. As individuals learn about the new level of market tightness, they slowly adjust their positions. Market tightness thus predicts price growth over multiple future periods. Guren (2018) models some fraction of sellers as using a backward-looking heuristic for house prices. This behavior generates momentum in house prices and corresponding inventory volatility, exacerbated by other, forward-looking households strategically timing their market entry to take advantage of the predictability of price growth.

[19] The unexplained portion of price growth may be attributed to supply and demand factors beyond the parsimonious objects "number of buyers" and "number of sellers" we study in this paper.

[20] For example, for-sale listings (supply) may have been depressed in 2023 due to owner-occupants feeling locked in by their pandemic-era low interest rates that are no longer available (Fonseca and Liu (2023)). Policies that unlocked existing homeowners (e.g. allowing mortgages to be portable as in some other countries) could significantly stimulate entry by joint buyer-sellers, but not necessarily by first-time home buyers.

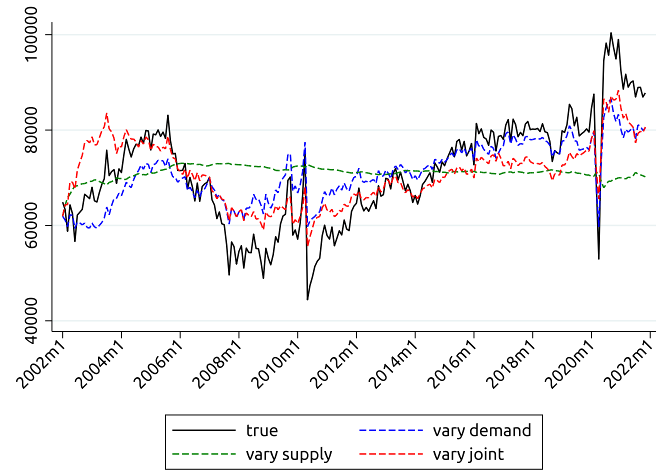

Fig. 7. Sales volumes with joint buyer-sellers.
Notes: "True" is the actual sales volume in the data. "Vary supply" is the counterfactual sales volume according to our model when the inflows of exclusive buyers and joint buyer-sellers are held fixed at their sample means, but the inflow of exclusive sellers varies as in the data. "Vary demand" is the counterfactual sales volume when the inflow of exclusive sellers and joint buyer-sellers are held fixed at their sample means, but the inflow of exclusive buyers varies as in our estimates. "Vary joint" is the counterfactual sales volume when the inflow of exclusive buyers and exclusive sellers are held fixed at their sample means, but the inflow of joint buyer-sellers varies as in our estimates.

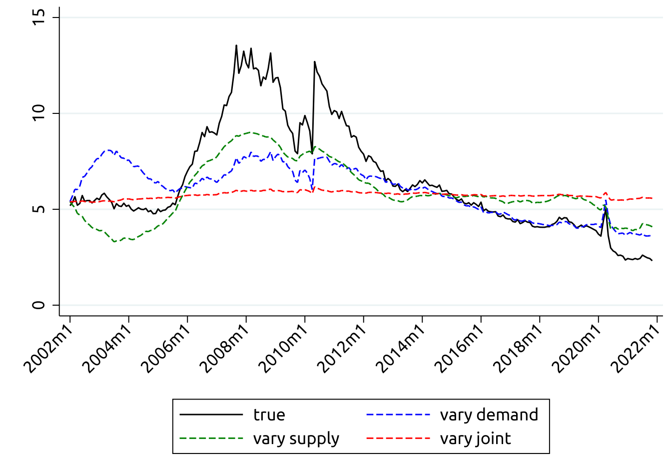

Fig. 8. Months' supply with joint buyer sellers.
Notes: "True" is the actual months' supply in the data. "Vary supply" is the counterfactual months' supply according to our model when the inflows of exclusive buyers and joint buyer-sellers are held fixed at their sample means, but the inflow of exclusive sellers varies as in the data. "Vary demand" is the counterfactual months' supply when the inflow of exclusive sellers and joint buyer-sellers are held fixed at their sample means, but the inflow of exclusive buyers varies as in our estimates. "Vary joint" is the counterfactual months' supply when the inflow of exclusive buyers and exclusive sellers are held fixed at their sample means, but the inflow of joint buyer-sellers varies as in our estimates.

In this section, we augment the model in Section 3 by adding a third type of household flowing into the market, joint buyer-sellers, or n^j. When a joint household enters, they increment the stock of both buyers and sellers, so Eq. (1) becomes:

$$s_{t+1} = s_t - s_t q_t^s(\theta_t) - s_t(1 - q_t^s(\theta_t))w^s + n_{t+1}^s + n_{t+1}^j \tag{9}$$

and Eq. (2) becomes:

$$b_{t+1} = b_t - b_t q_t^b(\theta_t) - b_t(1 - q_t^b(\theta_t))w^b + n_{t+1}^b + n_{t+1}^j \tag{10}$$

where n^s and n^b now refer to the inflow of households that are searching exclusively on one side of the market. The total inflow values, $n^s + n^j$ and $n^b + n^j$, are inferred using the same approach described in Section 4. We then estimate the fraction of home purchases that went to joint buyer-sellers (see the Appendix for details), and use this fraction along with the inferred inflows to separately estimate the time series of n^s, n^b, and n^j.

With the time series for market entry of the three types of households, we can repeat the analysis in Section 5 for this augmented model. We now simulate three counterfactual histories for each outcome variable, in which we hold two types of household inflows constant at their sample average while letting the third type vary as we estimate it did in reality.

The three counterfactual simulations of the number of home sales are shown in Fig. 7, along with the true data for comparison. As before, the simulation in which we vary only supply (in this case, sellers who are not joint buyer-sellers) explains none of the observed variation in total sales (the R^2 is negative). Simulations using variation in exclusive buyers or joint buyer-sellers both track the data much better. "Vary Demand" (varying the inflow of exclusive buyers) has an R^2 of 0.42, while "Vary Joint" has an R^2 of 0.44. Both the demand of exclusive buyers and joint buyer-sellers are important for determining overall sales volume. The supply provided by joint buyer-sellers will vary with their demand, but, as evidenced by the results in Section 5, this supply does not have any explanatory power on its own.

Turning to the results for months' supply, shown in Fig. 8, we see a very different pattern. Simulations using variation in the inflow of exclusive buyers ("vary demand") and of exclusive sellers ("vary supply") both have a moderate amount of explanatory power, with R^2 of 0.30 and 0.44, respectively. As we found in Section 5.2, the inflow of buyers and sellers both have independent effects on sale hazards. In contrast, the inflow of joint buyer-sellers ("vary joint") explains essentially none of the variation in the data, with an R^2 of 0.05. Joint buyer-sellers contibute one entrant to both the numerator and the denominator of market tightness, and so have little effect on sale hazards unless market tightness is very far from 1. To the extent that house price growth is tied to the months' supply of homes for sale, it is determined by the flows of exclusive buyers and sellers. The entry of joint buyer-sellers has close to no effect.

As can be seen by a comparison to the results in Section 5.2, demand from exclusive buyers explains only a fraction of the variation in months' supply relative to demand from all buyers. Even though the time series variation in the entry of joint buyer-sellers has no explanatory power by itself, these joint households are important for the simulated model's fit. This is because the presence of joint buyer-sellers dilutes the effect of exclusive-buyer or exclusive-seller entrants on overall market tightness. When the number of joint households on the market is low (as in the period from 2008 through 2011, for example), the marginal exclusive-buyer entrant has a relatively larger effect. By substituting the sample average inflow of joint buyer-sellers for their actual levels in the "vary demand" simulations in this section, the observed variation in exclusive buyers is not able to move market tightness by as much during this low demand period, resulting in a poorer fit to the data.

The results in this section extend but are consistent with those from Section 5. As in Section 5.1, we find that sales volume is determined almost entirely by the number of buyers rather than the number of sellers. Many of these buyers are also sellers, and our results in this section show that exclusive buyers and joint buyer-sellers are equally important for explaining volatility in sales volume. As in Section 5.2, we find that months' supply, and hence house price growth, are affected by both buyers and sellers, but in opposite directions (more buyers increases price growth whereas more sellers lowers it). Joint buyer-sellers have little effect on price growth because a joint household enters on both sides of the market and the two effects cancel out. For simplicity, we therefore use the baseline model with only two types of households throughout the remainder of the paper.

7. Contrast with a reduced form approach

An alternative, more reduced-form approach to partialling out the relative importance of supply and demand factors is to simulate various counterfactuals, varying or holding constant the terms in the accounting identity described in Eq. (1). This is the method used by Ngai and Sheedy (2020) (NS). In this section we replicate their motivating empirics, including a key element of their data construction. As described in Section 2, the aggregate NAR data NS use does not allow for measuring the number of new listings distinct from the number of withdrawals. We show why the reduced-form approach, along with the conflation of new listings and withdrawals, leads to opposite inferences about the importance of the supply side in determining sales volumes.

NS perform counterfactual simulations similar to those described in our Section 5, holding some determinants of sales fixed while allowing others to follow their observed path. The principal conceptual difference relative to our approach is that, rather than modeling the matching process and inferring the number of active buyers, they treat the sale hazard (q^s, in Eq. (1)) as a model primitive. Depending on the counterfactual being considered, this q^s is fixed as a constant or set following its observed historical path. The inflow of new listings, n^s, is treated similarly as in our approach. However, instead of using actual new listings and holding the withdrawal hazard constant as we do, given their aggregate data they must infer new listings as the monthly difference in for-sale inventories, net of sales. This construction is actually *net* new listings because it does not distinguish between an increase (decrease) in new listings and a decrease (increase) in withdrawals.

We follow their basic approach, once again initializing the market at the values observed in January 2002. We use Eq. (1) to simulate two counterfactual paths of active listings and sales. In the first, the sale hazard q^s is fixed at its sample mean, while the inflow of new listings n^s follows its observed path. To replicate their data construction, in this section we measure n^s as monthly new listings minus the monthly number of withdrawals and otherwise set the withdrawal hazard, w^s, equal to zero. In the second simulation, q^s is allowed to vary as it does in the data, while n^s is fixed at its sample mean. Simulated sales volumes under the two reduced-form scenarios, along with actual sales, are presented in Fig. 9. In contrast to the simulations using our full model (shown in Fig. 4), the reduced-form counterfactuals suggest the inflow of new listings has a powerful effect on the number of sales, and can explain much of the time-series variation. Unlike our simulated inflow of new buyers, variation in the sale hazard by itself does a poorer job explaining the number of sales. While there are some differences due to time period and additional data construction issues, these reduced-form findings qualitatively match those of NS.

Why does fixing the sale hazard as in NS lead to such different conclusions from the method we used in Sections 3 through 5? Recall from our model that the sale hazard, q^s, is a function of supply as well as demand: the more houses there are for sale, the more actively searching buyers are needed to maintain a particular sale hazard. This modeling of the sale hazard is consistent with the negative empirical correlation between for-sale inventory and the sale hazard shown in Fig. 1. Given the substantial observed time variation in the supply of new for-sale listings, a substantial amount of variation in the number of active buyers would be necessary to have kept the sale hazard fixed at a constant. Implicitly, the NS simulations in which listings vary but sale hazard is fixed involve considerable variation in demand.

This can be seen by taking the counterfactual sales data from the reduced-form simulations with varying supply and fixed sale hazard, and backing out the implied time series of demand (b_t) using the model described in Section 3. We apply the method described in Section 4.3 to this simulated data, and present the imputed level of counterfactual demand in Fig. 10. For comparison, we show the inferred time series of b_t based on our estimates from Section 4. As can be seen, the two series are quite similar. Given the actual inflow of new listings, a fixed sales hazard implies a time series of demand that follows the true historical demand series closely, with a small delay. The variation in sales volumes in NS's "vary supply" counterfactual mostly comes from variation in demand, even though the sales hazard is fixed.

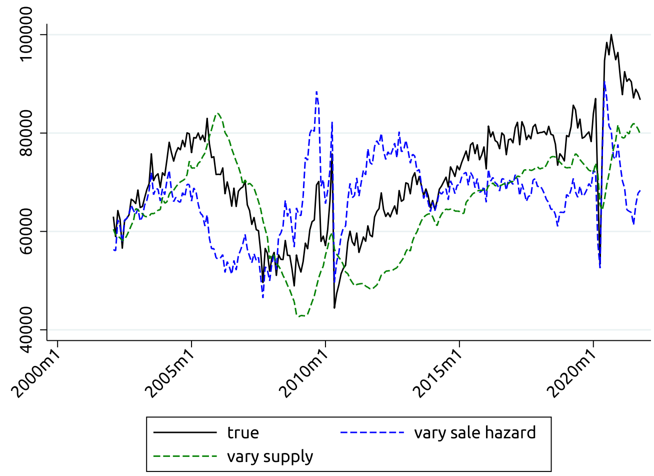

Fig. 9. Sales volume, reduced form.
Notes: "True" is the actual sales volume in the data. "Vary supply" is the counterfactual sales volume when the sale hazard is held fixed at its sample mean, but supply varies as in the data. "Vary sale hazard" is the counterfactual sales volume when supply is held fixed at its sample mean, but the sale hazard varies as in the data.

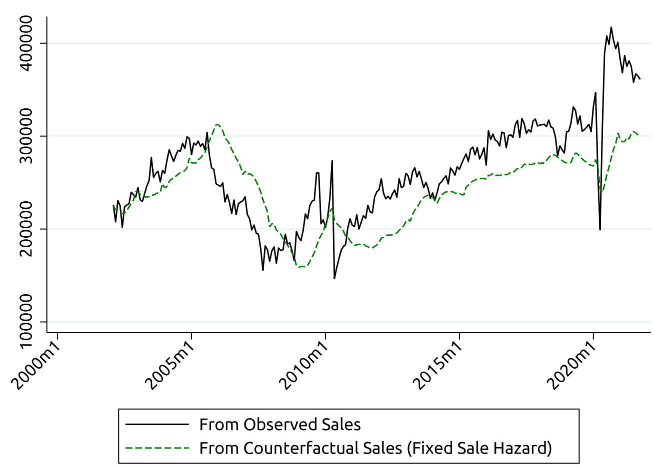

Fig. 10. Estimated number of active buyers.
Note: "From Observed Sales" shows the estimated number of active buyers implied by our model and the observed time series of listings and sales. "From Counterfactual Sales (Fixed Sale Hazard)" shows the number of active buyers implied by our model using a counterfactual sales volume series generated when the sale hazard is held fixed at its sample mean, but supply varies as in the data.

Including withdrawals in the "vary-supply" counterfactual contributes to the tight fit between the simulated and true data. As can be seen in Fig. 1, for example, during the years 2006–2008 new listings were still coming on the market at an elevated pace while sales volumes and sale hazard rates were falling and the housing boom turned to bust. Yet, the reduced-form "vary-supply" counterfactual sales volume takes a downturn at almost the same time as the true data do (see Fig. 9). This is possible because as shown in Appendix Figure 13, withdrawal rates rose as the sale hazard fell, growing about 40 percent from 2005 to their peak in 2008. The surge in withdrawals (and an unchanging sale hazard) depletes the counterfactual stock of for-sale listings faster than even the elevated level of new listings could replenish it, causing sales volumes in this simulation to fall as well. The combination of the reduced-form approach (effectively allowing zero supply-side crowd-out) and this conflation of withdrawals and new listings allows the

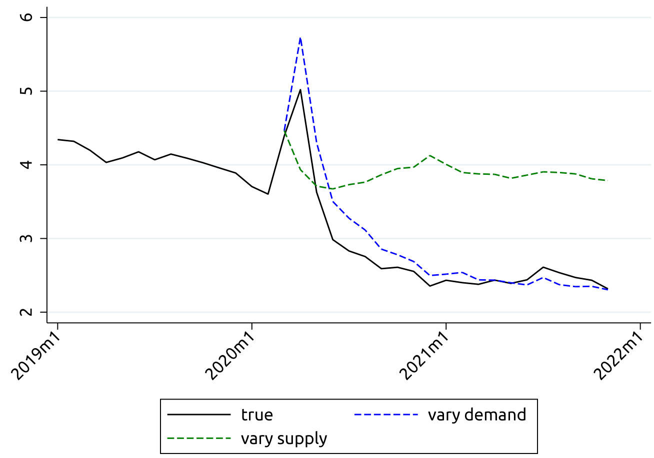

Fig. 11. Months' supply during COVID-19, observed and counterfactual.
Notes: Months' supply is equal to the inverse of the monthly sale hazard rate. "True" is the actual months supply in the data. "Vary supply" is the counterfactual months' supply according to our model when demand is held fixed at pre-pandemic (2019) levels, but supply varies as in the data. "Vary demand" is the counterfactual months' supply according to our model when supply is held fixed at pre-pandemic (2019) levels, but demand varies as in the data.

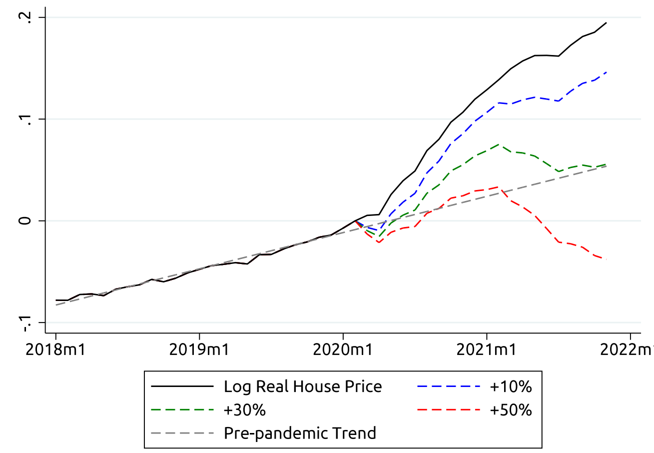

Fig. 12. Real house price during COVID-19, observed and counterfactual.
Notes: Shows log real house price under counterfactual supply. "+X" is a counterfactual where supply is set at a multiplier, X, of pre-pandemic (2019) supply levels for each month from March 2020 onward.

"vary-supply" simulations in this section and NS to fit the true sales data so well.[21]

The results of this section highlight the importance of taking market tightness, and its implications for the matching process, into consideration when evaluating the relative roles of demand and supply. The reduced-form results would suggest the supply of new listings, rather than demand for homes, is the most important factor in determining sales volumes. Our full set of results suggest that the opposite is true.

8. Implications

8.1. COVID-19 housing boom

Fig. 2 shows that during the COVID-19 pandemic, the housing market tightened considerably. After a brief dip at the onset of the pandemic, the sale hazard rate surged to record levels and house price growth also moved up to record highs. In this section, we use our model to decompose the tightening of the housing market during the pandemic into supply or demand factors.

A priori, the recent observed tightening in the housing market could be due to reduced supply or increased demand, or both. On the demand side, lower interest rates and widespread telework may have induced more buyers into the market. On the supply side, homeowners could be reluctant to list their home for sale during a pandemic, which could have reduced the for-sale supply. Generous mortgage forbearance programs and the foreclosure moratorium may also have reduced supply. Indeed, new listings plummeted at the onset of the pandemic.

Fig. 11 shows counterfactual months supply using our model under (i) fixed demand and true supply and (ii) true demand and fixed supply. When demand or supply is fixed, we set it at average 2019 (pre-pandemic) levels. At the very beginning of the pandemic, the vary-supply simulation drops below the true months' supply while the vary-demand simulation rises above the true months' supply, showing that some of the initial decrease in months' supply is driven by a decrease in new listings. As the pandemic progresses, however, the figure shows that stronger demand overtakes lower supply as the main factor behind the observed decrease in months' supply. By the middle of 2021, the contribution of reduced supply has disappeared and higher demand can explain essentially all of the decrease in months' supply since March 2020. We conclude that, outside of a brief shock at the beginning of the pandemic, reduction of supply was a minor factor relative to an increased number of buyers in explaining the tightening of housing markets.

We can also use our model to estimate how much additional supply would be needed to keep house prices on their pre-pandemic trend, given the observed increase in demand. Fig. 12 shows counterfactual house prices in which demand (n_t^b) is set at its actual estimated levels, but supply (n_t^s) is set at some multiplier, x, of average 2019 (pre-pandemic) levels. We assume the relationship between house price growth and months' supply during the pandemic is the same as the one shown in Fig. 2. We find that a value of $x = 1.3$ or greater is necessary to bring the counterfactual months' supply (and hence house price growth) back to its pre-pandemic trend by November 2021. This means that a 30% increase in the monthly number of homes coming on to the market would have been necessary to keep up with the pandemic-era surge in the number of buyers. This is a very large increase in supply. Since new construction typically accounts for about 15% of supply, our estimates imply that new construction would have had to increase by roughly 300% to absorb the pandemic-era surge in demand. One implication of this result is that policies targeted at increasing supply, for example construction subsidies or zoning reforms, would have done little to cool the pandemic house price boom in the short-run.

8.2. Interest rate elasticity

This section compares the sensitivity of the demand for and supply of homes to changes in interest rates, which is an important channel through which policy makers can influence the housing market.

We estimate the regression:

$$y_t - y_{t-12} = \alpha_0 + \alpha_1(frm_t - frm_{t-12}) + \epsilon_t \tag{11}$$

where y is the housing market variable of interest in month t and frm is the average monthly 30-year fixed mortgage rate in percentage points as reported in the Freddie Mac primary mortgage market survey. We estimate the regressions using our monthly sample between January 2002-November 2021.

Table 2 reports estimates of α_1 for different outcome variables and Appendix Figure 24 shows a binscatter presentation of the regression

[21] In addition to their headline counterfactuals, NS do attempt simulations that partial out the effects of withdrawals. However, lacking observations of individual listings, they are forced to assume an elasticity of withdrawals relative to sales hazard. This exercise does weaken the power of listings to explain sales in their paper, but does not eliminate it.

Table 2
Mortgage rate elasticity.

	Demand	Supply	Sales
30-yr Fixed Mortgage Rate	−8856.7***	2837.7*	−4292.2***
	(2185.4)	(1599.9)	(1450.0)
Observations	226	226	226

Standard errors in parentheses
* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.
Note: All variables are in 12-month changes. Demand is the demand for homes, as implied by the housing search model. Supply is new listings and Sales is sales volume. Newey–West standard errors with optimal lag-selection algorithm are shown.

results. The first column of the table shows that higher mortgage rates have a strong negative effect on buyer demand, n_t^b. A negative effect is expected because higher mortgage rates increase the cost of owning a home, which should decrease demand all else equal. A one percentage point increase in the mortgage rate is associated with about 9,000 fewer buyers entering the market in our sample counties. Relative to the sample average value of buyer demand, this is a decrease of about 10.4 percent, or a semi-elasticity of 10.4. Our estimates capture the extensive margin of changes in demand to changes in mortgage rates. Changes in mortgage rates can also affect demand through an intensive margin – for example, buyers choosing to purchase larger or smaller homes – and our estimates do not capture this margin.

Column 2 shows that home sales are also negatively associated with mortgage rates, but the magnitude of the effect is much smaller. The semi-elasticity of home sales to mortgage rates is estimated to be 6, about one-half the estimate of the demand semi-elasticity. Why are home sales much less mortgage rate sensitive than our estimate of buyer demand? One explanation is search frictions. Because it takes time for buyers to transact, home sales today reflect demand from a mix of periods in the past. This mixture effectively smooths the response of home sales to demand shocks, leading to attenuated estimates. A second reason is that, as column 3 shows, there is a small, positive association between mortgage rates and new listings.[22] Higher supply results in higher sales volume all else equal, so the negative relationship between demand and mortgage rates is somewhat offset by the positive relationship between supply and mortgage rates. An implication of these results is that the demand for homes is much more responsive to mortgage rates than simple regressions based on observables imply.

9. Conclusion

We use a housing search model to decompose fluctuations in home sales and prices into supply or demand factors. Simulations of the estimated model show that the number of buyers searching for homes drives short-run fluctuations in home sales and prices.

For longer-run changes in the housing market, supply may play a much larger role. For example, new supply today also increases supply in the future as today's buyer eventually sells her new home. Our simulations do not account for such a response as we are focused on the short run, but the accumulation of new supply (including new construction) likely explains more of the variation in sales volume over long horizons. Similarly, long-run levels of house prices may not be as closely connected to market tightness as the short-run price growth we consider in this paper. Understanding the relative importance of supply and demand and other factors for longer-run changes in the housing market remains a topic for future research.

[22] One potential explanation for a positive association between mortgage rates and new listings is investor activity. Investors may be more inclined to sell homes when interest rates rise and the expected returns from investing in assets other than housing become relatively more attractive.

Appendix A. Supplementary data

Supplementary material related to this article can be found online at https://doi.org/10.1016/j.jue.2023.103610.

References

Aladangady, Aditya, 2017. Housing wealth and consumption: Evidence from geographically-linked microdata. Amer. Econ. Rev. 107 (11), 3415–3446.
Albrecht, James, Gautier, Pieter A., Vroman, Susan, 2016. Directed search in the housing market. Rev. Econ. Dyn. 19, 218–231.
Anenberg, Elliot, Bayer, Patrick, 2020. Endogenous sources of volatility in housing markets: The joint Buyer–Seller problem. Internat. Econom. Rev. 61 (3), 1195–1228.
Anenberg, Elliot, Ringo, Daniel, 2021. Housing market tightness during COVID-19: Increased demand or reduced supply? Federal Reserve Note.
Anenberg, Elliot, Ringo, Daniel, 2022. The propagation of demand shocks through housing markets. Am. Econ. J.: Macroecon. 14 (3), 481–507.
Benmelech, Efraim, Guren, Adam, Melzer, Brian T., 2023. Making the house a home: The stimulative effect of home purchases on consumption and investment. Rev. Financ. Stud. 36 (1), 122–154.
Berger, David, Guerrieri, Veronica, Lorenzoni, Guido, Vavra, Joseph, 2017. House prices and consumer spending. Rev. Econom. Stud. 85 (3), 1502–1542.
Berger, David, Turner, Nicholas, Zwick, Eric, 2020. Stimulating housing markets. J. Finance 75 (1), 277–321.
Best, Michael Carlos, Kleven, Henrik Jacobsen, 2017. Housing market responses to transaction taxes: Evidence from notches and stimulus in the UK. Rev. Econom. Stud. 85 (1), 157–193.
Bhutta, Neil, Ringo, Daniel, 2021. The effect of interest rates on home buying: Evidence from a shock to mortgage insurance premiums. J. Monetary Econ. 118, 195–211.
Caplin, Andrew, Leahy, John, 2011. Trading frictions and house price dynamics. J. Money Credit Bank. 43, 283–303.
Carrillo, Paul E., Williams, Benjamin, 2019. The repeat time-on-the-market index. J. Urban Econ. 112, 33–49.
Carrillo, Paul E., de Wit, Eric R., Larson, William, 2015. Can tightness in the housing market help predict subsequent home price appreciation? Evidence from the United States and the Netherlands. Real Estate Econ. 43 (3), 609–651.
DeFusco, Anthony A., Nathanson, Charles G., Zwick, Eric, 2022. Speculative dynamics of prices and volume. J. Financ. Econ. 146 (1), 205–229.
DeFusco, Anthony A., Paciorek, Andrew, 2017. The interest rate elasticity of mortgage demand: Evidence from bunching at the conforming loan limit. Am. Econ. J.: Econ. Policy 9 (1), 210–240.
Diaz, Antonia, Jerez, Belen, 2013. House prices, sales, and time on the market: A search-theoretic framework. Internat. Econom. Rev. 54 (3), 837–872.
Fonseca, Julia, Liu, Lu, 2023. Mortgage Lock-In, Mobility, and Labor Reallocation. Working Paper.
Gabrovski, Miroslav, Ortego-Marti, Victor, 2019. The cyclical behavior of the Beveridge Curve in the housing market. J. Econom. Theory 181, 361–381.
Gabrovski, Miroslav, Ortego-Marti, Victor, 2021. On the Slope of the Beveridge Curve in the Housing Market. Working Paper.
Genesove, David, Han, Lu, 2012. Search and matching in the housing market. J. Urban Econ. 72 (1), 31–45.
Grindaker, Morten, Karapetyan, Artashes, Moen, Espen R, Nenov, Plamen, 2021. Transaction Sequencing and House Price Pressures. CEPR Discussion Paper No. DP16351.
Guren, Adam M., 2018. House price momentum and strategic complementarity. J. Polit. Econ. 126 (3), 1172–1218.
Guren, Adam M, McKay, Alisdair, Nakamura, Emi, Steinsson, Jón, 2021. Housing wealth effects: The long view. Rev. Econom. Stud. 88 (2), 669–707.
Guren, Adam M., McQuade, Timothy J., 2020. How do foreclosures exacerbate housing downturns? Rev. Econom. Stud. 87 (3), 1331–1364.
Gyourko, Joseph, Molloy, Raven, 2015. Regulation and housing supply. In: Handbook of Regional and Urban Economics, vol. 5, Elsevier, pp. 1289–1337.
Han, Lu, Strange, William C., 2015. The microstructure of housing markets: Search, bargaining, and brokerage. In: Handbook of Regional and Urban Economics, vol. 5, Elsevier, pp. 813–886.
Hanushek, Eric A., Quigley, John M., 1980. What is the price elasticity of housing demand? Rev. Econ. Stat. 449–454.
Head, Allen, Lloyd-Ellis, Huw, Sun, Hongfei, 2014. Search, liquidity, and the dynamics of house prices and construction. Amer. Econ. Rev. 104 (4), 1172–1210.
Hedlund, Aaron, 2016. Illiquidity and its discontents: Trading delays and foreclosures in the housing market. J. Monetary Econ. 83, 1–13.
Karahan, Fatih, Rhee, Serena, 2019. Geographic reallocation and unemployment during the great recession: The role of the housing bust. J. Econom. Dynam. Control 100, 47–69.
Krainer, John, 2001. A theory of liquidity in residential real estate markets. J. Urban Econ. 49 (1), 32–53.
Mian, Atif, Rao, Kamalesh, Sufi, Amir, 2013. Household Balance Sheets, Consumption, and the Economic Slump*. Q. J. Econ. 128 (4), 1687–1726.

Mian, Atif, Sufi, Amir, 2011. House prices, home equity-based borrowing, and the US household leverage crisis. Amer. Econ. Rev. 101 (5), 2132–2156.

Molloy, Raven, Nathanson, Charles G., Paciorek, Andrew, 2022. Housing supply and affordability: Evidence from rents, housing consumption and household location. J. Urban Econ. 129, 103427.

National Assocation of Realtors, 2019. 2019 Profile of Home Buyers and Sellers. Technical report.

Ngai, L. Rachel, Sheedy, Kevin D., 2020. The decision to move house and aggregate housing-market dynamics. J. Eur. Econom. Assoc. 18 (5), 2487–2531.

Novy-Marx, Robert, 2009. Hot and cold markets. Real Estate Econ. 37 (1), 1–22.

Ortalo-Magne, Francois, Rady, Sven, 2006. Housing market dynamics: On the contribution of income shocks and credit constraints. Rev. Econom. Stud. 73 (2), 459–485.

Piazzesi, Monika, Schneider, Martin, 2009. Momentum traders in the housing market: Survey evidence and a search model. Amer. Econ. Rev. 99 (2), 406–411.

Ringo, Daniel, 2022. Monetary policy and home buying inequality. Available at SSRN 4220206.

Saiz, Albert, 2010. The geographic determinants of housing supply. Q. J. Econ. 125 (3), 1253–1296.

Smith, Eric, 2020. High and low activity spells in housing markets. Rev. Econ. Dyn. 36, 1–28.

Zabel, Jeffrey E., 2004. The demand for housing services. J. Housing Econ 13 (1), 16–35.